Methodology v2.3.2 · effective 2026-08-21 · document revised 2026-08-31 · rendered 2026-09-08 from ccir-v2-notebooks@255323a · ccir.io/documents/methodology

CCIR Reference Rates — Public Methodology

Version: 2.0 · **Status:** External methodology (published). This is the outward-facing framework, rules, and governance for the CCIR reference rates, written for users, counterparties, and IOSCO alignment. It deliberately does **not** publish the full source panel or the collection playbook; those are maintained in the internal methodology and available to a regulator or auditor under appropriate terms. Representative sources are named for transparency; the complete roster is proprietary so the benchmark cannot be trivially replicated or gamed.

1. Overview & scope

- CCIR publishes daily **reference rates for GPU compute** — productivity-adjusted $/GPU-hour marks used for credit agreements, collateral valuation, and structured finance.
- **Silicon scope:** NVIDIA data-center accelerators — Hopper, Blackwell, and headline Ampere SKUs — indexed by chip × tier. The framework is silicon-agnostic; additional accelerator families are admitted to a published series only once they clear the same panel and gate requirements as the incumbent silicon.
- **Geographic scope:** a curated US/EU headline; other regions surface in their own regional cells.
- **Rate basis:** published **list-asks** (public rate-card offers), not executed transactions — the same posted-price basis used by established price-reporting agencies.

2. The Intelligence Factory Taxonomy (tier framework)

CCIR grades every offer by the *deployment environment*, not just the chip. The premium ladder:

- **T1IF — Hyperscale Intelligence Factory:** top-grade production cluster fabric (high-bandwidth InfiniBand or RoCE-class interconnect) **and** a rich managed bundle (managed control plane, enterprise SLA, observability) **and** dedicated hosts.
- **T2IF — Datacenter Intelligence Factory:** a mid-grade-or-better disclosed cluster fabric (InfiniBand / RoCE / EFA) **and** dedicated hosts.
- **T3IF — Marketplace Intelligence Factory:** shared or co-located hosts, **or** dedicated hosts without a disclosed production-grade fabric — marketplace, sub-host slices, best-effort.

The tier is set by the deployment environment along three axes — cluster fabric, host tenancy, and managed bundle depth. The exact grading cutoffs are applied consistently across the panel and maintained in the internal methodology. The same chip can sit in different tiers at different providers; the tier prices the **fabric + host + bundle**, which is what a lender's collateral is actually exposed to. A parallel source-class taxonomy carries long-tail and inference SKUs.

3. Data sourcing

- **Principle:** a **curated, multi-source panel** of public rate cards spanning the market structure — hyperscalers (e.g., **Amazon Web Services**), institutional neoclouds (e.g., **CoreWeave**), specialized and regional providers (e.g., **Scaleway**), and distributed marketplaces.
- **Panel roster:** the full membership and per-source configuration are proprietary and confidential — disclosing them would enable gaming of the benchmark and copying of the collection process. Representative members are named above for transparency.
- **Inclusion criteria:** first-party operators publishing observable $/GPU-hour list-asks in a stable currency. Excluded: aggregators and resellers (double-count), token-denominated venues (FX-contaminated), wrapped-unit products (credits/tokens), and quote-only offers.
- **Collection methods, SKU mappings, and endpoints** are maintained in the internal methodology.

4. Aggregation

- **Operator-equal median:** each named source counts **once** — the headline is the median of per-source medians, so a provider publishing many SKUs cannot dominate. No single quote moves the mark.
- **Horizon:** a freshness-first fallback ladder (same-day preferred, widening to a short trailing window only as needed to meet the gate).
- **Publication gates:** a series publishes only above a minimum distinct-source count and minimum observation count; emerging regional cells may publish as **Provisional** and auto-promote. Below the floor, a series is computed but withheld (**Shadow**).
- **Composition stability:** a frozen-panel and last-good carry-forward mechanism prevents a transient collection gap from moving the headline as a composition artifact.

5. Publication & confidence

- Statuses: **Published · Provisional · Shadow.**
- Confidence tiers are keyed on source breadth.
- The headline is the **median**; mean, inter-quartile range, and extremes are published as diagnostics.

6. Chip & term scope

- The headline is the **on-demand** reference rate. Committed-term (reserved) pricing is captured and maintained in the internal methodology; if and when it is surfaced publicly it will publish as a

breakout, never blended into the on-demand headline.

7. Governance & IOSCO

- **Frozen reference panel:** panel membership is fixed and versioned; changes are governance events, logged and dated. (The roster itself is confidential.)
- **Methodology events:** every rule or scope change is recorded with a before/after rationale.
- **Independence:** CCIR publishes reference rates and does not operate a trading venue against them.
- **IOSCO alignment:** the benchmark's rules, governance, and controls are disclosed here; panel composition is held confidential to prevent gaming — consistent with accepted price-reporting-agency practice.

8. What is not published here (by design)

The full source roster, per-source tier assignments, exact SKU and endpoint mappings, collection cadence, the precise grading cutoffs, and the operational collection process. These are maintained in the internal methodology and available to a regulator or auditor under appropriate terms, but are not published so the benchmark cannot be trivially replicated or gamed.